



Faculty of Engineering and Applied Science

Department of Electrical, computer and Software Engineering

SOFE 3650 – Software Design & Architecture

Deliverable 2: ADD Iteration 1 & 2

Due Date: November 19th, 2025

Course Instructors: Dr Hani Sami

CRN: 45740 - Group #8

Student Name	Student Number
Christopher Lui	100912564
Geraline Chavez	100890130
Nicholas Furtado	100908880

1. The Design Process

The design process for the AI-Powered Digital Assistant Platform (AIDAP) is guided by the Attribute-Driven Design (ADD) method. This process marks the transition from requirements analysis to architectural reasoning, where the system's functional and quality drivers are translated into design decisions. The goal of this phase is to establish an initial architectural vision that captures the most influential drivers identified in Phase 1, particularly performance, scalability, and maintainability, and to outline how the architecture will evolve through successive iterations. Subsequent steps will progressively refine this high-level structure into concrete components, interactions, and deployment views.

1.1 ADD Step 1: Review Inputs

The purpose of this step is to review the main design inputs from Phase 1, which serve as the foundation for the first architectural iteration. These inputs include the functional and quality requirements, constraints, and architectural concerns that will guide all subsequent design decisions.

Design Inputs Summary

Category	Details
Design Purpose	To design a cloud-based AI assistant platform (AIDAP) that centralizes academic and administrative information into a single intelligent interface for students, lecturers, and administrators. The goal is to enhance accessibility, ensure seamless integration between institutional systems (LMS, email, and calendar), and deliver rapid, reliable responses through an adaptive and scalable architecture.
Primary Functional Requirements	UC-1 – Merge and Display Multiple Academic Calendars: Enables integration of different university systems (LMS, email, calendar). This requirement drives the need for an Integration Layer and API Gateway to unify data sources in real time. It impacts interoperability and maintainability.

	<p>UC-4 – Perform Server Maintenance with Zero Downtime: Ensures the system remains available while updates or patches are applied, ensuring minimal downtime. This requirement drives the use of redundant microservices, rolling deployments, and container orchestration (Kubernetes), supporting availability and maintainability.</p> <p>UC-5 – Handle High Traffic Efficiently: Supports up to 5,000 concurrent users during peak load (e.g., exams). It drives the choice of a cloud-native microservices architecture and load-balancing mechanisms to achieve performance and scalability.</p>
--	---

Quality Attributes

Scenario ID	Quality Attribute	Importance to Customer	Difficulty of Implementation	Reason
QA-1	Performance	High	High	Essential for fast responses under 5,000 users; drives the need for efficient Django views and API Gateway optimization.
QA-2	Availability	High	Medium	Required for continuous service during maintenance; relies on redundancy and failover mechanisms.
QA-3	Usability	High	Medium	Important for natural user interaction (text/voice); moderate effort via

				interface design and AI training.
QA-4	Privacy and Security	High	High	Mandatory institutional policy enforcement; complex because of access-control and data-encryption requirements.
QA-5	Scalability	High	High	Needed for traffic spikes; container orchestration and load-balancing add complexity.
QA-6	Maintainability	Medium	Low	Facilitated by modular microservices; lower complexity due to Django's MVT structure.
QA-7	Adaptability	Medium	Medium	Enhances personalization; requires machine-learning feedback loops.

From this list, QA-1 (Performance), QA-2 (Availability), QA-5 (Scalability), and QA-6 (Maintainability) are selected as *architectural drivers* because they most strongly influence responsiveness, uptime, scalability, and maintainability, which align with AIDAP's cloud-native goals.

Constraints	All of the constraints (CNST-1 to CNST-6) discussed in Phase 1 are included as drivers.
Concerns	All architectural concerns (CRN-1 to CRN-8) from Phase 1 are included as drivers.

1.2 Iteration 1: Establishing an Overall System Structure

This section presents the results of the activities performed in the first iteration of the ADD process.

The goal of this iteration is to establish an initial system-level structure for the AI-Powered Digital Assistant Platform (AIDAP) based on the most influential architectural drivers identified in the previous step.

1.2.1 Step 2: Establish the Iteration Goal by Selecting Drivers

This first iteration of the Attribute-Driven Design process is guided by the architectural goal of establishing AIDAP's overall system structure.

The focus is on defining a modular, scalable, and fault-tolerant architecture that ensures high performance and continuous availability, while supporting seamless integration between internal and external university systems.

Although this iteration is primarily driven by the architectural concern CRN-5: Maintain a clearly defined and logically structured architecture, the architect must also account for all the quality, constraint, and concern drivers influencing the system's initial design.

The primary design drivers for this iteration are summarized below:

Architectural Drivers Considered

Type	Driver ID	Influence on Design
Quality Attribute	QA-1 – Performance	The system must handle 5,000 concurrent users and respond to queries within 2 seconds. This drives the need for load balancing, efficient API endpoints, and asynchronous request handling.
Quality Attribute	QA-2 – Availability	Continuous operation must be maintained during maintenance and system updates. This leads to redundant deployment, container orchestration, and failover support.

Quality Attribute	QA-5 – Scalability	The architecture must dynamically allocate resources during high-traffic periods, requiring containerization (Docker/Kubernetes) and cloud-native infrastructure.
Quality Attribute	QA-6 – Maintainability	Modular microservices and isolated deployment pipelines must be implemented to simplify updates and reduce downtime.
Constraint	CNST-1	The system must handle 5,000 concurrent users while maintaining a maximum response time of 2 seconds.
Constraint	CNST-6	The system must employ a cloud-native architecture for scalability and portability across environments.
Architectural Concern	CRN-5	The architecture must remain logically structured and modular to support long-term scalability and ease of maintenance.

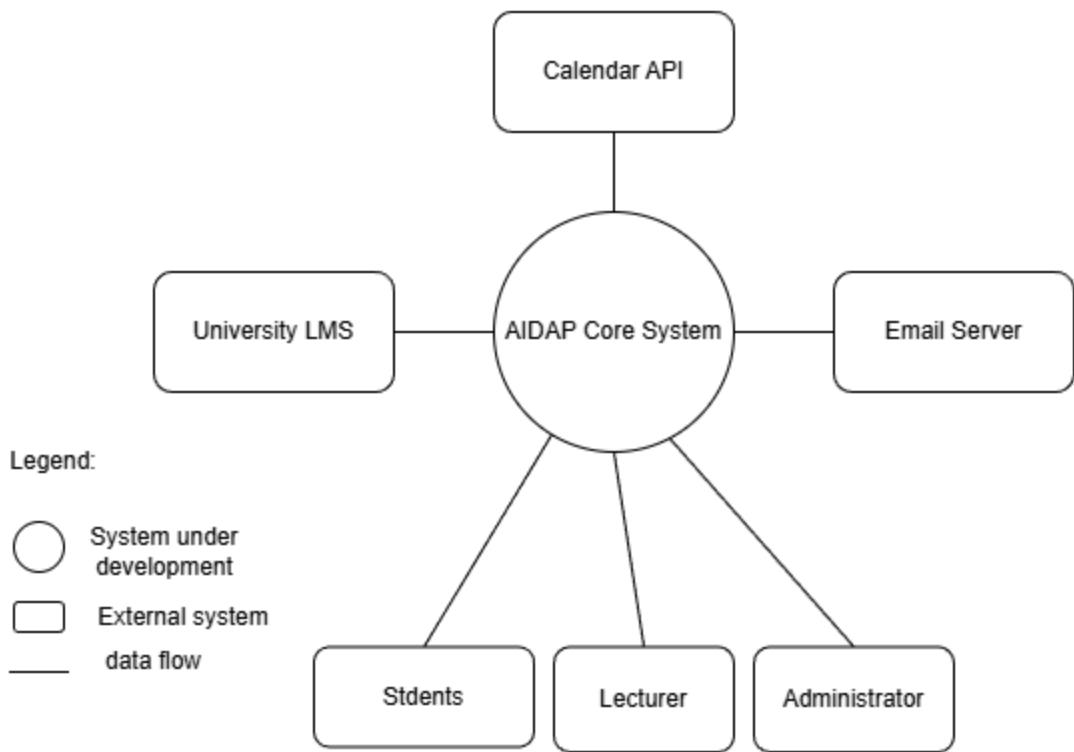


Figure 1– Context Diagram for the AIDAP System

This diagram shows the external entities (users and connected systems) interacting with the AIDAP Core System. Lines represent data exchange, consistent with the ADD context view format.

1.2.2 Step 3: Choose One or More Elements of the System to Refine

This is a greenfield development effort; therefore, in this iteration, the element to refine is the entire AIDAP Core System. The purpose of this step is to define the system boundaries and identify its external entities, which include institutional systems and user roles.

The Context Diagram (Figure 1) illustrates these relationships, showing how AIDAP interacts with connected systems, such as the University LMS, Calendar API, and Email Server, as well as end-users, Students, Lecturers, and Administrators. This context definition serves as the foundation for further architectural decomposition in subsequent iterations.

1.2.3 Step 4: Choose One or More Design Concepts That Satisfy the Selected Drivers

In this iteration, the selected design concepts define the high-level structure of the AIDAP system. They are based on the main architectural drivers such as scalability, availability, and maintainability. These decisions establish the foundation for system decomposition and guide future design refinements.

Design Decisions & Locations	Rationale
Establish AIDAP core system as primary element of refinement	Defining the system boundary is essential before decomposition (CRN-5)
Define external entities	(External Entities include: Students, lecturers, administrators, LMS, registration system, Calendar API, Email Server, etc) Identification of actors defines how AIDAP must interact with external institutional systems. This supports maintainability of the system (QA-6).
Adopt a cloud-native architecture for the Core System	Required by CNST-6. Enables dynamic scaling, and better system availability. It ensures the platform can support high amount of concurrent users (5000) and continuous operations (QA - 2, QA - 5)
Use containerization for all deployable components	Supports portable deployment and efficient scaling with an architecture that is cloud native. Directly contributes to performance and maintainability (QA - 1, QA - 6)
Structure the system using modular microservices	Allows for independent scaling and deployment while also ensuring faults are isolated. It is vital for maintainability and keeps the architecture modular (QA - 6, CRN - 5).
Introduce API gateway as access point for interactions	Centralize authentication, routing, and monitoring. Improves performance, simplifying the process of external integration. It also allows for a single entrypoint across the system services (QA - 1).
Require processing using message queues	This reduces latency and ensures that the system meets the response constraints. It also enhances fault tolerance and supports system scalability. (CNST-1, QA-2, QA-5).
Define the systems initial high level functional domains	Establishes the structure needed for iteration 2 decomposition. These domains include the Chat UI, AI Engine Domain, Integration Domain, and User/Profile Domain. This ensures the system architecture remains logically structured, modular, and maintainable (CRN-5, QA-6).

1.2.4 Step 5: Instantiate Architectural Elements, Allocate Responsibilities and Define Interfaces

The instantiation design decisions considered and made are summarized in the following table:

Design Decisions &	Rationale
-------------------------------	------------------

Locations	
Create initial domain model	Identifies the main AIDAP entities (User, Academic Calendar, LMS Data, Sessions, Messages). This clarifies the foundation of the system and speeds up later design steps.
Identify use cases inside domain objects	Maps UC-1 (Merge Calendars), UC-4 (Zero-downtime Maintenance), UC-5 (High-Traffic Handling) to the domain model to ensure all key behaviors are supported.
Decompose domain objects into functional domains	Splits the system into clear areas (Chat UI, AI Engine, Integration, User/Profile). Helps distribute workload evenly and supports modular microservice design.
Instantiate API Gateway module	Provides a single controlled entry point for all users and external university systems. Handles routing, authentication, and monitoring. Supports performance and maintainability.
Create core microservice modules	Instantiates containerized modules for Chat UI (Django), AI Engine, Integration Service, and User/Profile Service. The User/Profile Service handles authentication, user data, and authorization logic.
Introduce Message Queue for async tasks	Supports notification delivery, background processing, and traffic spikes. Ensures fast responses during peak load (5000+ users).
Allocate storage components	Defines an SQL RDBMS for structured data and Object Storage for large files/archives. Ensures reliable data persistence and backup handling.
Define preliminary REST interfaces	Establishes the basic communication paths between the API Gateway and each microservice. Detailed interfaces will be refined in Iteration 2.

The results of these instantiation decisions are recorded in the next step.

1.2.5 Step 6: Sketch Views and Record Design Decisions

This is the modified version of the Rich Client architecture based on the decision made in step 5.

Figure 2- Rich Client architecture diagram

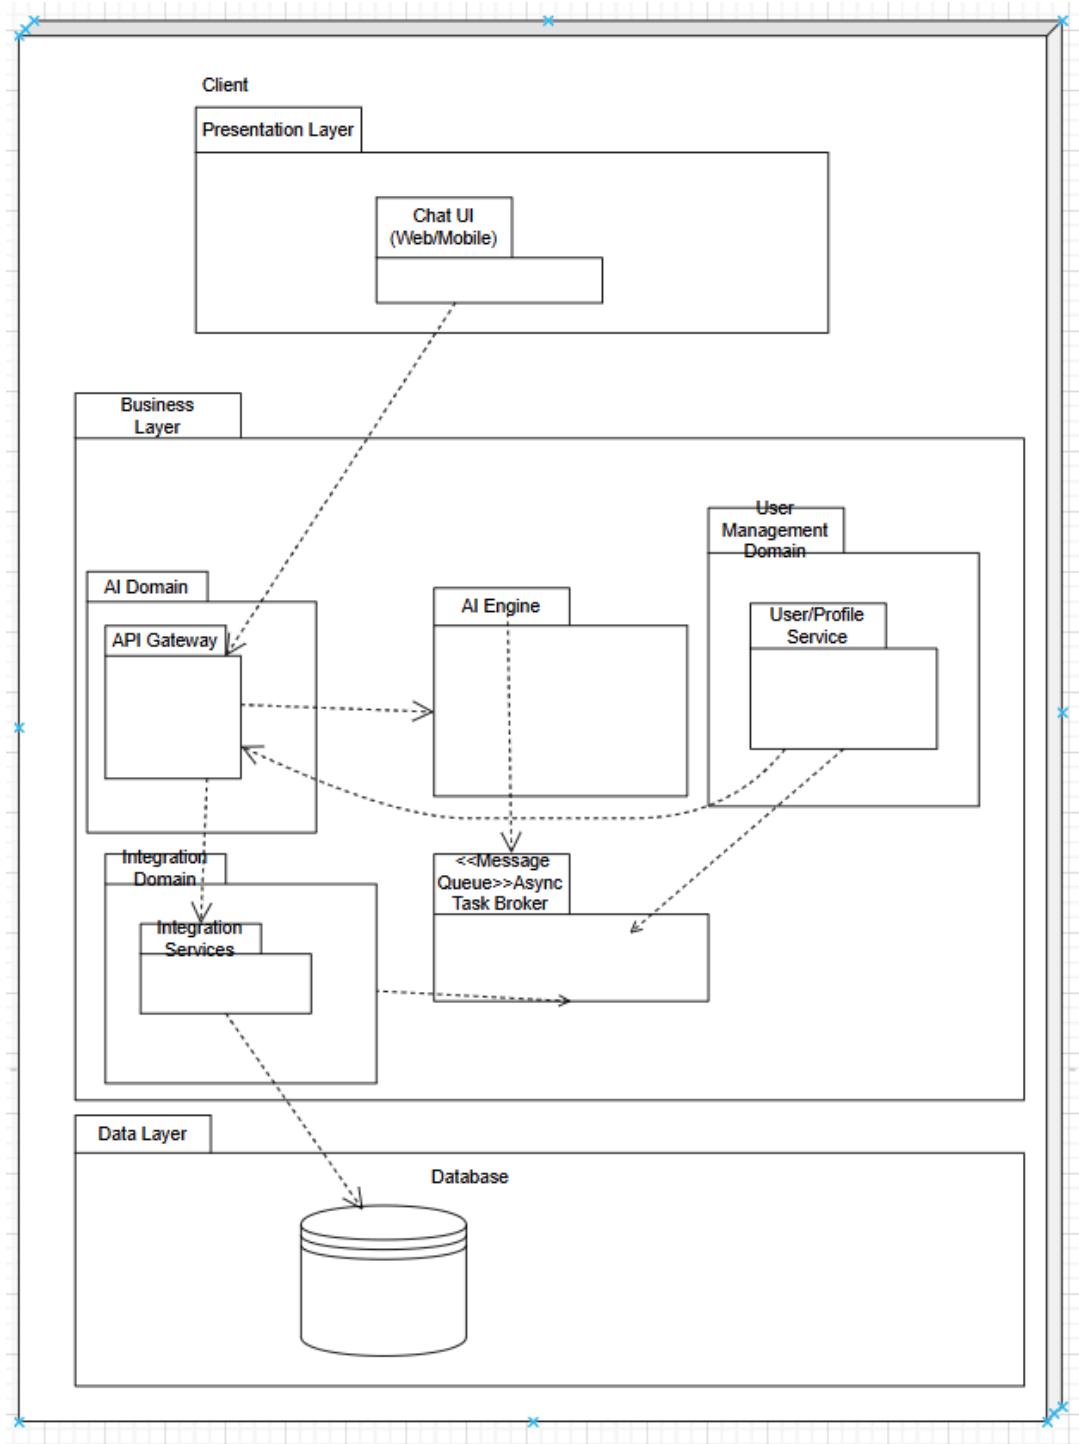
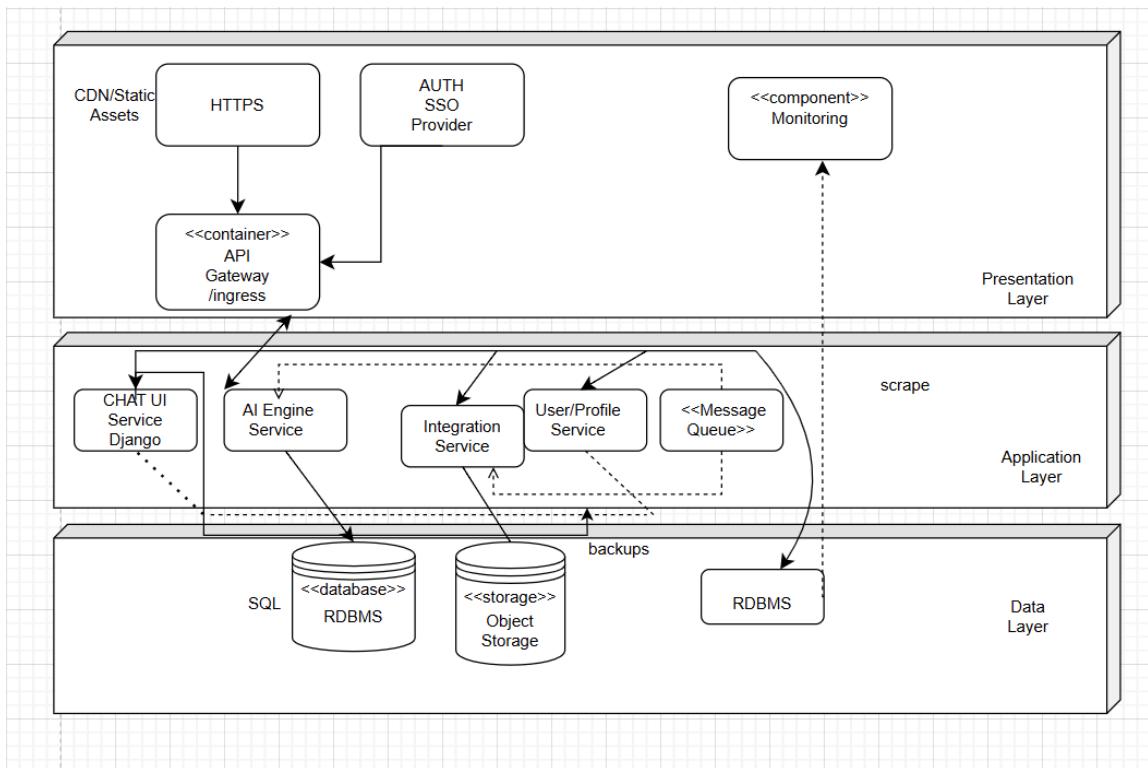


Table based on the elements in the rich client architecture

Element	Responsibility
Client Side	Represents the user-facing environment where the system is accessed through web or mobile platforms. It is responsible for capturing user input and displaying the system's responses in an intuitive and interactive way.
Chat UI (Web/Mobile)	Provides the main interaction interface for users. It allows them to communicate with the system through natural language and forwards queries to the backend via the API Gateway.
API Gateway	Serves as the central access point for all client-side requests. It manages request routing, authentication, and communication between the client and the AI Engine. This design simplifies external access and improves system security and scalability.
AI Engine	Represents the core processing component of the system. It analyzes user requests, applies intelligent logic or machine learning models, and generates responses based on system data and integrated services.
User/Profile Service	Manages user authentication, authorization, profile data, and access control policies. Ensures secure user interactions with the system.
Message Queue	Handles asynchronous communication between services by decoupling time-sensitive tasks such as notifications and background processing from real-time interactions.
Integration Services	Handles communication with external systems and internal services, such as academic databases, scheduling services, or analytics modules. It

	ensures seamless data exchange between the AI Engine and other subsystems.
Database	Responsible for storing and managing all persistent system data, including user records, interaction history, and domain-specific information. It supports data consistency, retrieval, and long-term storage.

Figure 3- Deployment Diagram



Responsibilities of the elements in the deployment diagram

Element	Responsibility

CDN / Static Assets	Distributes static frontend resources such as CSS, images, and scripts to reduce server load and improve performance.
HTTPS Interface	Enables secure communication between clients and the system using encrypted protocols.
AUTH SSO Provider	Provides authentication and single sign-on services for users.
API Gateway (Ingress)	Acts as the main entry point for client requests, handling routing and security enforcement.
Chat UI Service (Django)	Manages user interaction and sends requests to backend services.
AI Engine Service	Processes user queries using AI models and generates intelligent responses.
Integration Service	Manages communication with external systems and internal services.
User/Profile Service	Handles user account management and profile-related information.

Message Queue (RabbitMQ / Kafka)	Enables asynchronous communication between services to support background tasks and high traffic processing.
RDBMS (SQL Database)	Stores structured system data such as user records, logs, and system information.
Object Storage	Stores large files, system assets, and backup data.
Monitoring Component	Collects system logs, performance metrics, and service health information for monitoring and diagnostics.

Relationships between some elements in the deployment diagram

Relationship	Description
Client → API Gateway	Client requests are routed to the system through the API Gateway over HTTPS.
API Gateway → AI Engine Service	The API Gateway forwards user queries to the AI Engine for processing.

API Gateway → User/Profile Service	The API Gateway routes authentication and user-related requests to the User/Profile Service.
AI Engine → Message Queue	AI tasks are offloaded to the Message Queue for asynchronous execution.
Integration Service → Message Queue	Integration tasks are handled asynchronously using the Message Queue.
Core Services → RDBMS	All main services store and retrieve structured data from the relational database.
Core Services → Object Storage	Services use object storage for large files and backup handling.
System Services → Monitoring Component	All services send logs and performance metrics to the Monitoring system.
Object Storage → RDBMS (Backups)	Database backups are stored in the object storage system.

1.2.6 Step 7: Perform Analysis of Current Design and Review Iteration Goal and Achievement of Design Purpose

Not Addressed	Partially Addressed	Completely Addressed	Design Decisions Made During the Iteration
---------------	---------------------	----------------------	--

	UC-1		API Gateway selected to support calendar merging requests efficiently.
	UC-4		Microservice decomposition prepared for zero-downtime updates.
	UC-5		Cloud-native, containerized services chosen to support high-traffic load.
		QA-1	Message Queue introduced to improve performance under 5,000+ concurrent users.
		QA-2	Redundant, container-based deployment strategy introduced for availability.
		QA-5	Modular microservices allow independent scaling during peak load.
		QA-6	Clear functional domains and service separation support maintainability.
	CNST-1		High-traffic constraint addressed through asynchronous processing and efficient routing.
		CNST-6	Cloud-native architecture enforced using containers and scalable storage.
	CRN-5		Architecture structured into presentation, application, and data layers.

1.3 Iteration 2: Refined Architectural Design

Iteration 2 represents the more refined and detailed design decisions we have made prior to development of the system. This step will allow us to form a coherent strategy before beginning the project, which will streamline the development process and prevent overlap.

1.3.1. Step 2. Establish Iteration Goal by Selecting Driver

The goal is to identify specific structures that aid functionality. Clarifying the use cases to be focused on in this iteration allows us to focus on the task at hand and not get sidetracked. In iteration 2, we decided to focus on the use cases we did not take into consideration in iteration 1.

Primary Use Cases	Description
UC-2: Maintaining Student Privacy	A student requests to view their friend's grade for the recent midterm. The AI model detects that the user making the request is not authorized to view the material they are asking for, so it politely, but firmly, declines.
UC-3: Detecting Drop in Participation	The AI model notifies the lecturer that participation with assignments has decreased significantly over the last month. The lecturer requests to view the average GPA of his class and sees that it has lowered. They make an announcement regarding extra credit for bonus assignments.
UC-6: Referencing Chat History	The student uses the assistant to help with studying for an exam over the course of a week. The assistant is able to pick up where they left off the previous session by analyzing the previous messages in the chat. The assistant is also able to explain concepts in a way that the student can understand, based on their chat history.

1.3.2. Step 3. Choose one or more elements of the system to refine

In this iteration, the main element to refine is the self-learning capabilities of the AI model. Since the purpose of this assistant is to guide students and lecturers through their academic careers with as little friction as possible, it is essential that the model is able to train effectively over numerous user interactions.

Another important element to refine is the security and privacy of users when interacting with the model. We must ensure that sensitive information such as grades are not given away by the

assistant, even when directly prompted. To accomplish this, we will need a way to direct data flow to the correct sections of the system, to allow for careful filtering of outputs from the assistant.

1.3.3 Step 4. Choose one or more design concepts that satisfy the selected drivers

This step highlights the decision concepts used in the planning process. Each of these concepts were carefully selected to ensure the most efficient initial modelling as possible. We made our decisions based on the amount of detail relative to the time it would take to complete each step.

Design Decision	Explanation
Initial Domain Model	A functioning domain model is crucial to the development of any system, even more so for ones pertaining to AI. Without identifying clear objects and entities within the system, development will be chaotic and random, contributing to working time and costs. Also important to distinguish differences between classes since privacy is a large concern.
Mapping of domain objects to use cases	Organizing the entities, operations and data flow of each use case is essential to ensuring all architectural drivers have been accounted for. An effective alternative to using mapping tables would be a traceability matrix, although this method could become difficult to manage with many drivers to consider.
Decompose domain objects into aggregates and value objects	Breaking the large entities within each domain object into smaller components lessens the possibility of leaking private data, as well as simplifies the implementation of frameworks later on in development. The alternative would be to skip decomposition entirely, but this is not recommended as domain objects may be unmanageable and tightly coupled.
Organizing system into layered architecture	Using the layered architecture style ensures a clear separation of concerns, which is important due to the private nature of certain objects. Isolating sensitive information from

	<p>the machine-learning algorithms will also allow for more organized workflow.</p> <p>An alternative we considered was blackboard architecture, as it is highly effective at analyzing data and integration, which could be a solid choice considering UC-3. Besides that, however, blackboard architecture is not typically used for the basic chat retrieval aspects of the system.</p>
Implementing Django and React frameworks	<p>Our back-end framework of choice is Django, not only because of its familiarity, but also due to its numerous built-in features such as user authentication. React can create dynamic graphs and analytics for our front-end, as well as interactive UI elements for grades and alerts.</p> <p>Another alternative we considered was flask, for its simplicity in integrating AI models, but has far fewer built in functions compared to Django.</p>
Data Transfer Objects	<p>Combines with layered architecture to provide a structured method for transferring data between layers. This enforces privacy by monitoring what information is shared with certain layers.</p> <p>The alternative to DTOs would simply be to directly use domain models, but we decided against it due to it being less efficient overall.</p>

1.3.4. Step 5. Instantiate Architectural Elements, Allocate Responsibilities, and Define Interfaces

The instantiation design decisions made in this iteration are summarized in the following table:

Design Decision	Explanation
Domain instantiation	Each entity within the system must be identified and modeled clearly. This must be done thoroughly to ensure that all drivers are included within the design. Only an initial model is required, so that we are able to begin

	development as soon as possible.
Map drivers to concrete services	Dedicated service classes should be implemented for each use case identified. This will prevent tight coupling, as well as simplify development while maintaining privacy of data.
Decompose domain objects into layer-specific modules	Each use case must be accounted for and satisfied within the modules. The decomposition of domain objects will allow for thorough inspection of each module, ensuring that no use case has been neglected. This task is best suited for multiple team members, as to make sure nothing is left unchecked.
Implement DTOs to define boundaries between layers	Introduce DTOs for each layer to control the data that flows between boundaries. Creating separate DTOs for lecturers, students and other entities guarantee that sensitive information is kept within its own boundary. This also allows for simplified API implementation and looser coupling.
Connect Components using Django	Connect architectural elements using chosen frameworks, allowing it to become runtime executable. Django uses app configuration to register modules, URL routing to map views to endpoints, and APIViews to expose services.
Create browser-based interface using React	Conceptual screens and data flows are implemented as functional react components. UI components are wired to backend contracts and services. React also enforces privacy and access routes in the client side.

1.3.5. Step 6 - Sketch Views and Record Design Decisions

As a result of the decisions made in step 5, several diagrams are created.

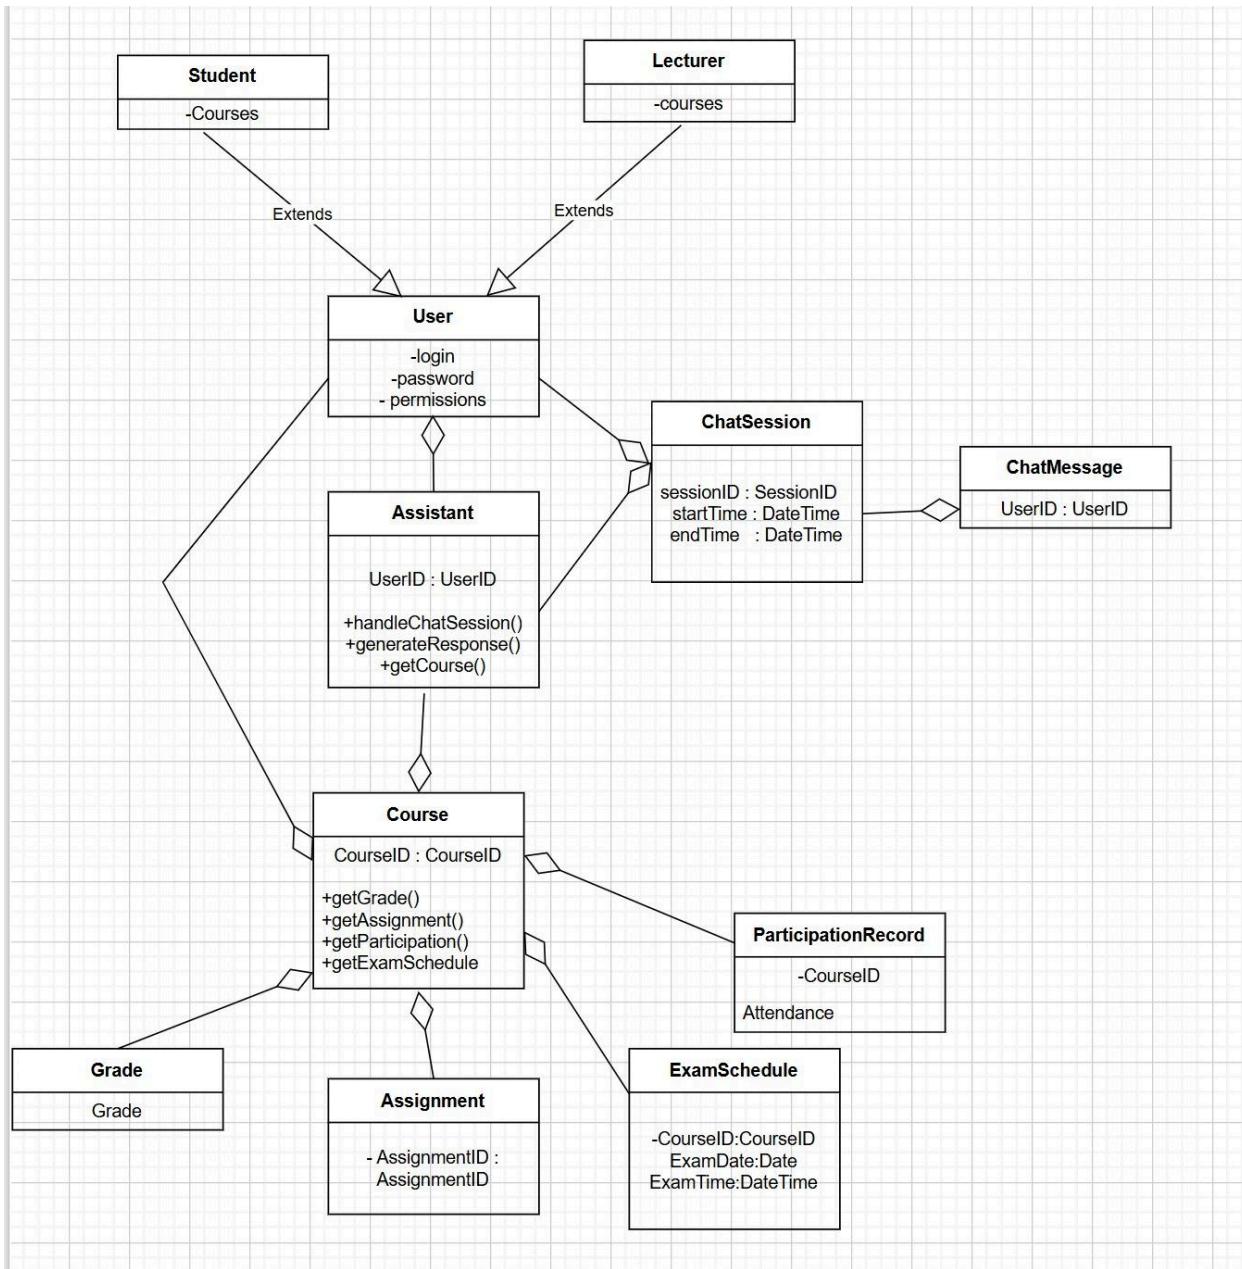


Figure 4 - Initial Domain Model

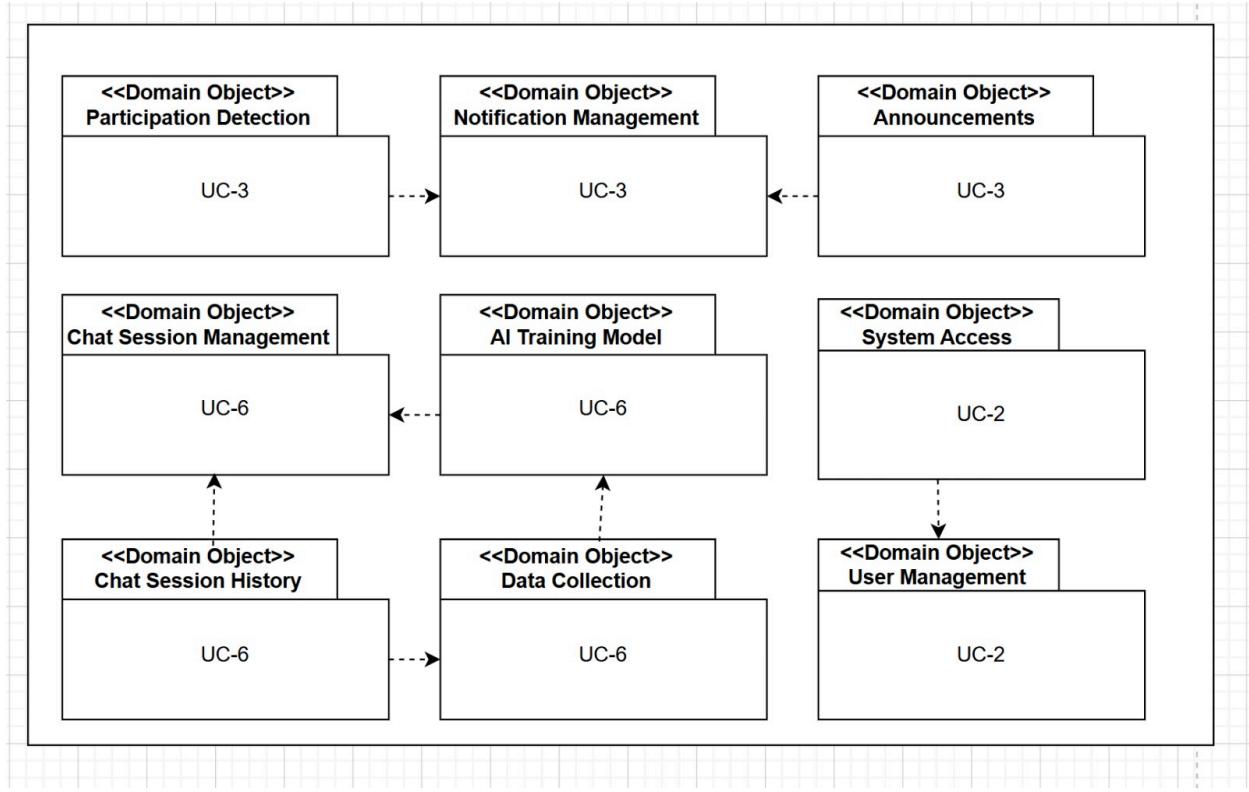


Figure 5 - Domain objects associated with the use case

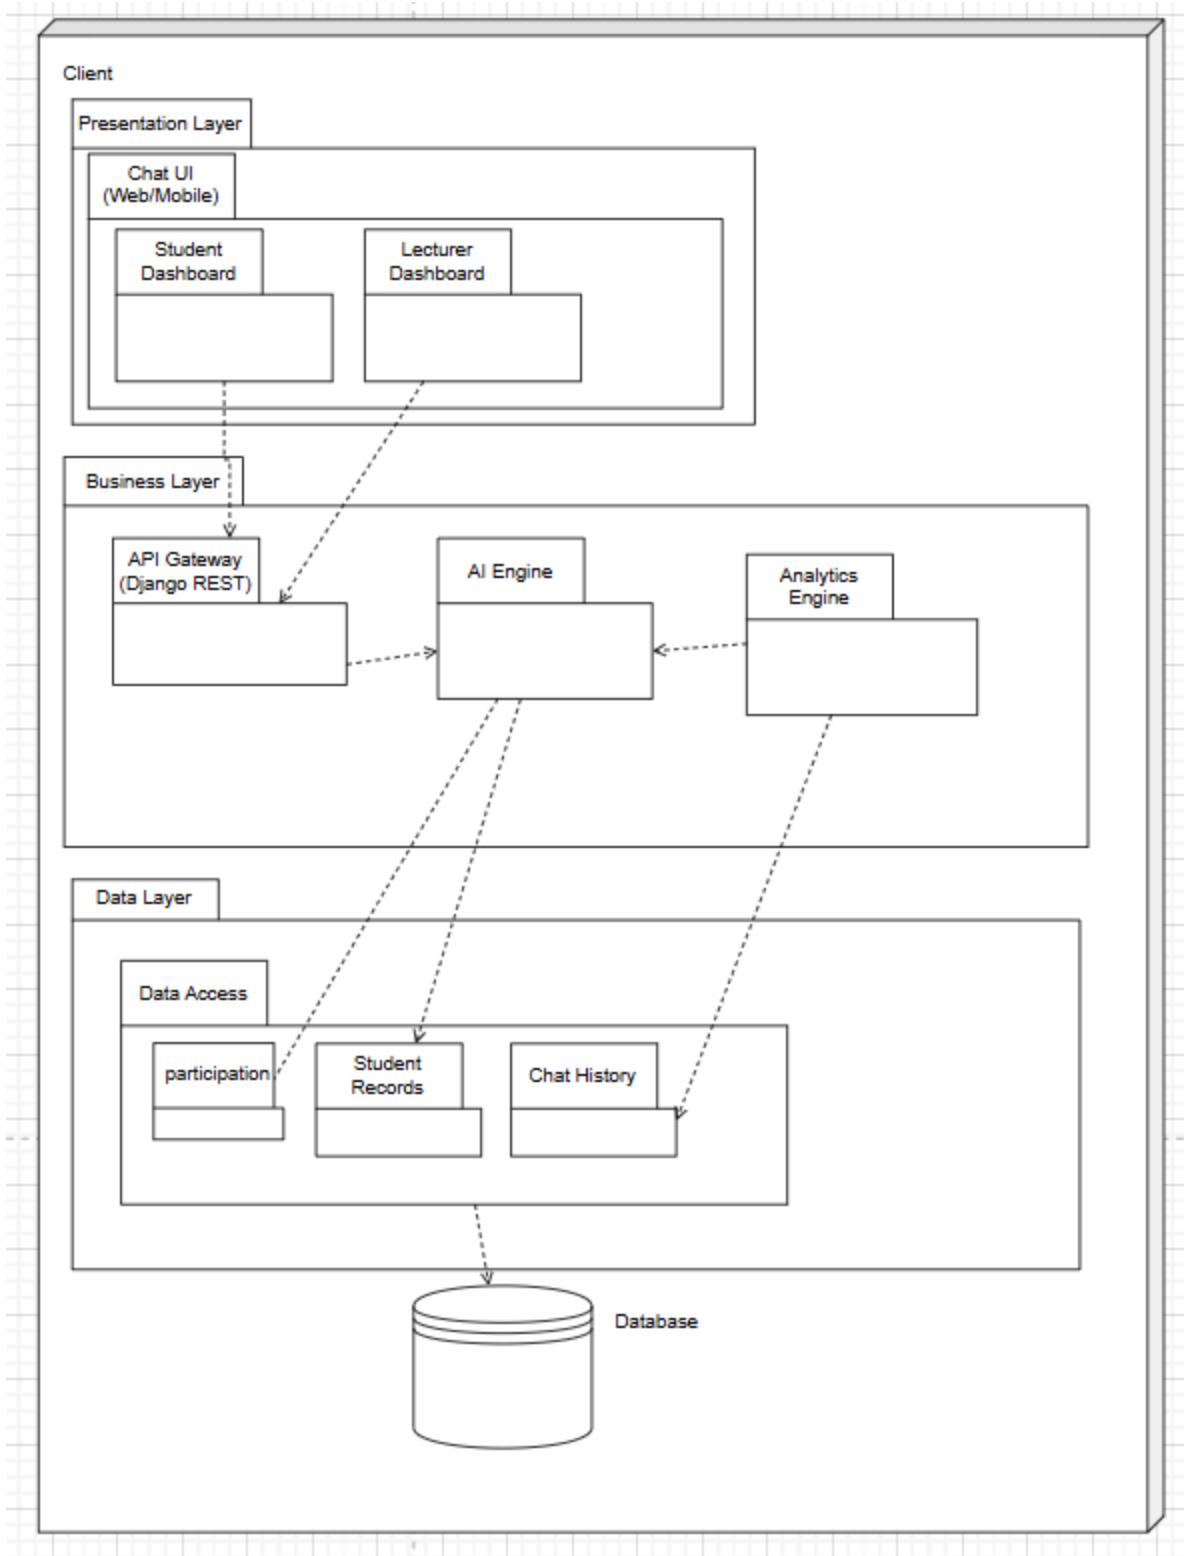


Figure 6 - Modules that support the primary use cases

Element	Responsibility
Chat UI (Web/Mobile)	User interface for the assistant, available on the website and on mobile
Student Dashboard	Student view of the dashboard
Lecturer Dashboard	Lecturer view of the dashboard
API Gateway (Django REST)	Allows users to interact with the assistant
AI Engine	Uses machine learning to find patterns in data to make decisions
Analytics Engine	Processes data using queries and statistical methods
Participation	Contains data about attendance and submissions during lectures
Student Records	Contains data concerning grades for all students
Chat History	Contains previous conversations with the assistant to train the model

UC - 2: Maintaining Student Privacy

Figure # shows an initial sequence diagram for UC - 2 (Maintaining student privacy). It shows how the system checks user credentials prior to accessing and providing sensitive information. Upon the user's request, the AIDAP system verifies student credentials and passes the information to the record system, where it then confirms that the requested information is not equal to the account information making the request. The AIDAP system then sends a polite prompt to the user letting them know that they are unauthorized to access the requested information.

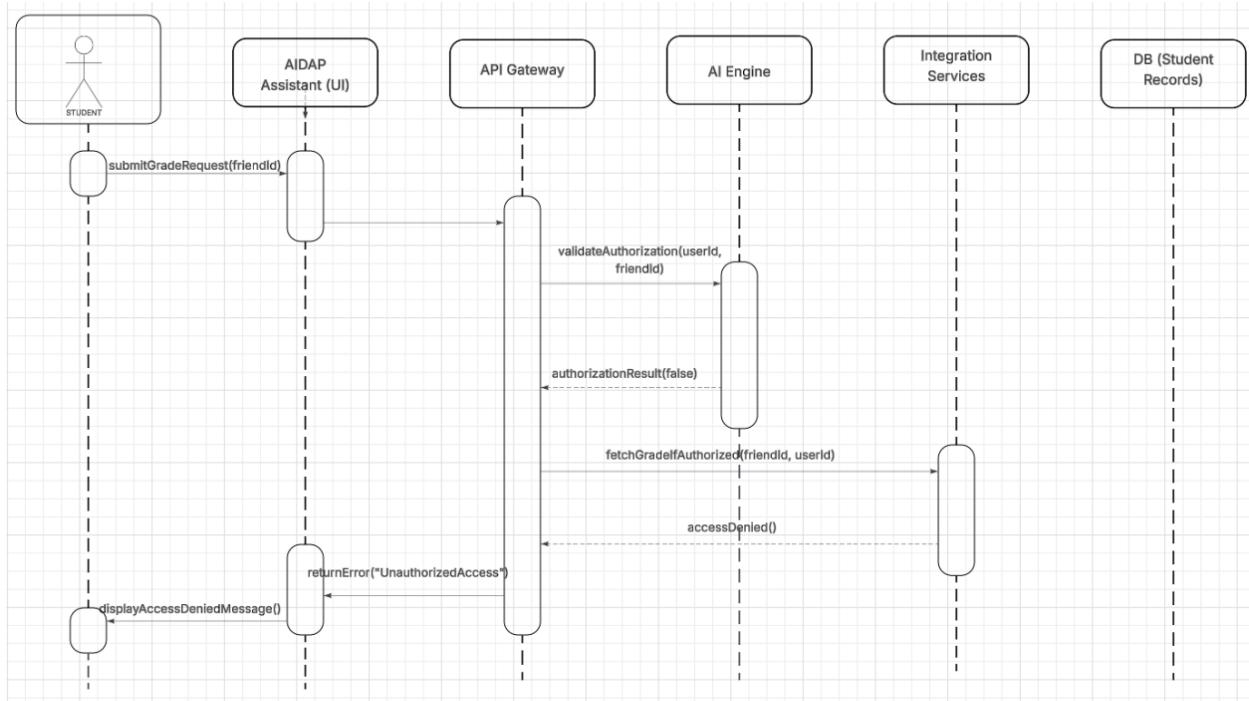


Figure 7 Sequence Diagram for UC - 2

From the interactions presented and identified in the sequence diagram above, initial methods for interfaces of all the interacting elements present can be identified:

Method Name	Description
Element: AIDAP Assistant (UI)	
<code>submitGradeRequest (friendId)</code>	Student submits request to view their friends grades
<code>displayAccessDeniedMessage ()</code>	User receives a prompt stating their access is denied
Element: API Gateway	
<code>returnError ("UnauthorizedAccess")</code>	Returns authorization failure
Element: AI Engine	
<code>validateAuthorization (userId, friendId)</code>	Compares and determines if logged in account and requested account are equal or vary
Element: Integration Services	

<code>fetchGradeIfAuthorized(friendId, userId)</code>	Attempts to fetch grade if accounts are equal, if not it denies access
<code>accessDenied()</code>	Replies with a denied access

UC - 3: Detecting Drop in Participation

Figure # shows an initial sequence diagram for UC - 3 (Detecting Drop in Participation). It shows how the system periodically monitors user participation within courses and provides lecturers with notifications on its observations. The analytics tracker built into the AIDAP system tracks the user participation across courses and when a drop is present, it generates a notification and prompts the lecturer of its observations. The lecturer is then able to access GPA information from the students record system from which they are able to use the AIDAP system to generate and publish a message to all of the users registered within the course.

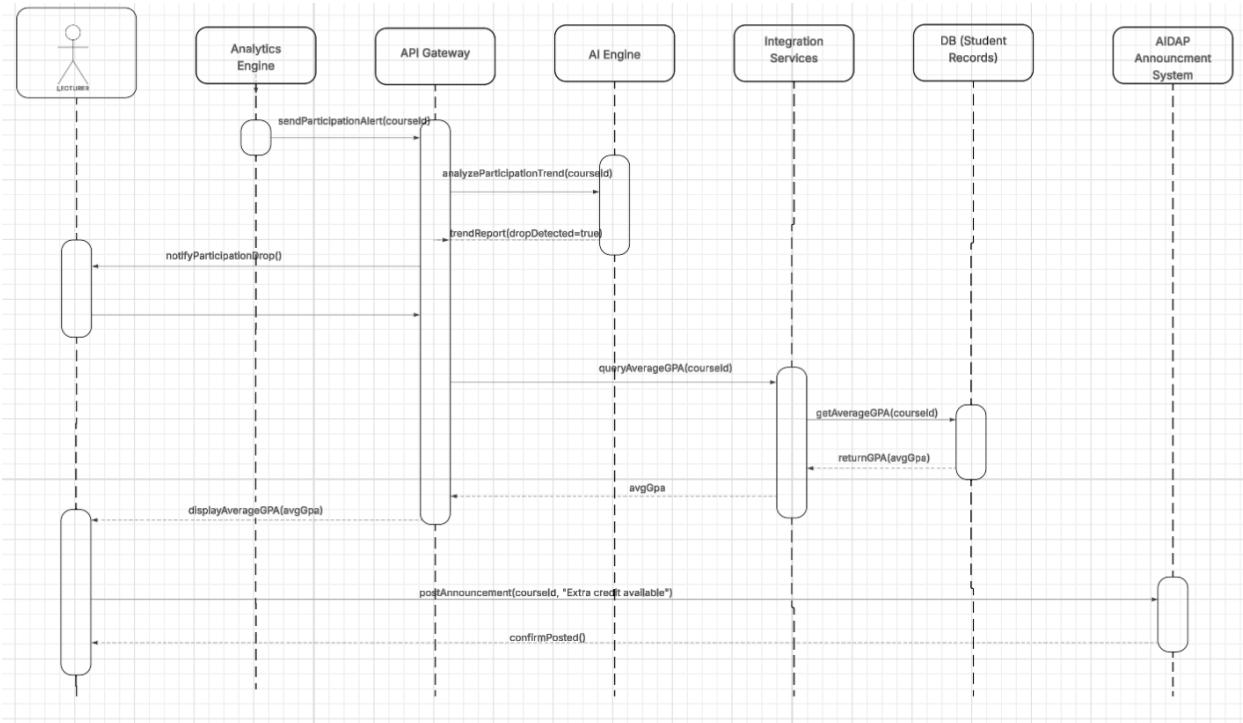


Figure 8 Sequence Diagram for UC - 3

From the interactions presented and identified in the sequence diagram above, initial methods for interfaces of all the interacting elements present can be identified:

Method Name	Description
-------------	-------------

Element: Analytics Engine	
sendParticipationAlert (courseId)	Sends an alert that participation within the course has dropped
Element: API Gateway	
notifyParticipationDrop ()	Notifies the lecturer of the drop in participation
Element: AI Engine	
analyzeParticipationTrend (courseId)	Confirms and analyzes the trends in the course participation
Element: Integration Services	
queryAverageGPA (courseId)	Retrieves the classes GPA
Element: DB (Student Records)	
getAverageGPA (courseId)	Executes GPA query
Element: AIDAP Announcement System	
postAnnouncement (courseId, message)	Posts extra credit announcement to the course registrants

UC - 6: Referencing Chat History

Figure # shows an initial sequence diagram for UC - 6 (Referencing Chat History).

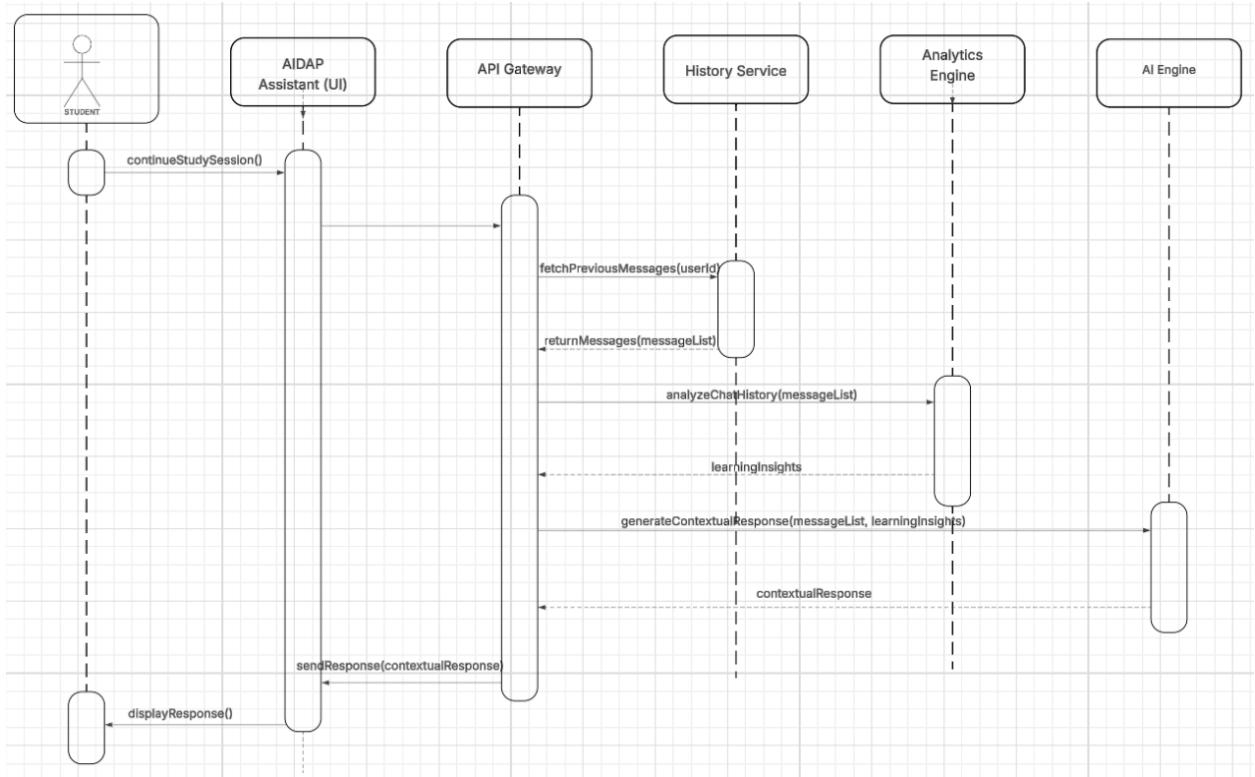


Figure 9 Sequence Diagram for UC - 2

From the interactions presented and identified in the sequence diagram above, initial methods for interfaces of all the interacting elements present can be identified:

Method Name	Description
Element: AIDAP Assistant (UI)	
<code>continueStudySession ()</code>	Student resumes a past study session
<code>displayResponse ()</code>	Displays response
Element: API Gateway	
<code>sendResponse (contextualResponse)</code>	Sends the final generated response to the display
Element: History Service	
<code>fetchPreviousMessages (userId)</code>	Retrieves the most recent message to continue the chat

Element: Analytics Engine	
analyzeChatHistory (messageList)	Analyzes past content for learning patterns
Element: AI Engine	
generateContextualResponse (messageList, learningInsights)	Generates the next set of responses based on past learning and analyzed patterns

1.3.6 Step 7 - Perform Analysis of Current Design and Review Iteration Goal and Achievement of Design Purpose

Not Addressed	Partially Addressed	Completely Addressed	Design Decisions Made During the Iteration
	UC - 2		Required modules that support Student privacy have been identified.
	UC - 3		Required modules that support analytics, alerts, and retrieval have been identified.
	UC - 6		Required modules supporting history retrieval and analysis have been identified.
QA - 1			No relevant decisions made, as it is necessary to identify the element that participate in the Use Case that is associated in the scenario
QA - 2			No relevant decisions made, as it is necessary to identify the element that participate in the Use Case that is associated in the scenario
	QA - 5		Scalability partially addressed through layered architecture.
	QA - 6		Maintainability partially addressed through modular services, and clear layering.
	CNST - 5		Performance partially supported through layered architecture and separation of concerns.

CNST - 6			No relevant decisions made, as it is necessary to identify the element that participate in the Use Case that is associated in the scenario
		CRN - 5	Logically structured architecture with clearly defined layers, and domain decompositions. Completely supports modularity and maintainability.